

SkillneX — Alignment Review

Prepared for: Product reviewer / stakeholder **Prepared by:** Engineering **Date:** 2026-04-24 **Subject:** Review of SkillneX_Before_After_Review.pdf against the MVP build **Companion docs in repo:** docs/office-hours-findings.md, docs/product-validation.md, PLAN.md, DESIGN.md

The document in hand

The Before-After Review PDF describes SkillneX as a **per-employee review deep dive** built on a **competency framework**:

- Each employee gets 5 competencies scored (Client Focus, Domain Application, Growth & Development, Operational Excellence, Perseverance & Impact)
- Each competency has **Manager Score vs Data-Derived Score** with a Gap and a status (OUTLIER / ALIGNED / REVIEW)
- Each score is backed by **4 verified signals** per competency, pulled from specific systems (Calendar, Slack, CRM, LMS, SmartSheet, Scribe, Jira, Salesforce, Rework Log, Project Log, Peer Feedback)
- Side-by-side layout: **Part A** shows the original manager one-sentence plus the employee's 5-bullet self-assessment (with flags: "generic", "self-reported, unverified"). **Part B** shows the SkillneX-generated score, the gap, and the narrative citing those 4 signals
- Final summary: three framed findings (Manager Comment Problem, Hidden Impact Problem, Calibration Risk)

The origin story is sharp: *"The employee wrote 5 evidence-backed bullet points per competency. The manager wrote one sentence. Both said Successful. Only one could defend it. SkillneX was built to fix that imbalance."*

What the MVP already gets right

Most of the product primitives the Before-After PDF depends on are already built and tested in the current MVP.

Before-After PDF feature	MVP status
Core thesis: verified data vs self-reported, outlier detection, narrative generation	Built in full
Positioning: replace the 3 days of manual data-gathering before the review, not replace Lattice	Already the homepage hero copy
Data-Derived Score vs existing rating, with gap flag	Implemented as the score-vs-rating-mismatch anomaly, labeled "Rating disagrees with data" in the UI
Ready-to-paste manager narrative, data-backed, strict-narration mode	NarrativeCard component with Copy button. Real Haiku 4.5 run: 94% pass rate against the safety guard, zero HR-action language leaks
Impact Profile framing (not "Employee ROI" in HR-facing UI)	Renamed in the Path B commit per product-validation findings

Before-After PDF feature	MVP status
Safety architecture: no invented numbers, no HR recommendations	7 safety rules in the cached system prompt, post-LLM guard rejects inventions, 64 unit tests green
Anomaly chip semantics (OUTLIER / ALIGNED / REVIEW)	Same concept, different labels: anomaly, success, neutral chips
Multi-source signal aggregation	Parser already reads multiple sheets and joins them into one employee record
Systems-connected tiles	/integrations page renders 4 tiles (demo-mocked)

Signal: the product philosophy, safety discipline, positioning, and most of the UI primitives are shared between what we built and what the PDF describes. We are not off-track.

Structural gaps vs the Before-After PDF

The PDF operates at a different zoom level than the MVP. Specifically:

Gap	What the PDF expects	What the MVP has
Competency model	5 named competencies per employee, each independently scored	One value score per employee, per department
Side-by-side review	Part A (original review: manager rating + one-sentence comment + employee 5-bullet self-assessment) alongside Part B (data-derived score + verified signals + narrative)	Single detail view with computed metrics only. No ingestion of the original review.
Verified signals per competency, labeled by source system	4 signals per competency (e.g. Calendar, Slack, CRM, Peer Mentions for Client Focus)	One combined signal blob per employee; sources not labeled
Manager-comment + employee-self-assessment ingestion	First-class inputs to the review	Not captured
Generic-comment detector	Flags one-sentence manager comments as non-specific	Not implemented
Per-competency narrative generation	5 narratives per employee, one per competency, each citing the signals for that competency	One narrative per employee
Three-findings summary	Cycle-close summary: Manager Comment Problem, Hidden Impact Problem, Calibration Risk	Not present
PDF export of the review	Export the per-employee Before-After layout as a PDF	Not built (the office-hours PDF pipeline can be adapted)

Gap	What the PDF expects	What the MVP has
Review Time Saved metric	Shown in the hero stat strip	Not measured

The honest read

The MVP answers the HR calibration question: *"Show me the whole team and flag who to investigate."* 110 employees ranked, outliers on top, a scatter plot for calibration meetings.

The Before-After PDF answers the review-writing question: *"Pre-fill this employee's review with data so the manager does not have to."* One employee, five competencies, side-by-side original-vs-data, verified signals, ready-to-paste prose per competency.

Both are legitimate. The PDF is the stronger product because it is closer to the actual workflow it replaces (a manager sitting down to write a review with a blank page). The team view is a secondary HR-admin workflow that complements it.

The MVP proves the thesis. The Before-After PDF describes the next zoom-in.

Reuse, not rewrite

Every primitive the Before-After PDF needs already exists in the code. Adding the new layer is additive, not destructive.

- **Gap detection between computed and existing rating** — reused as the per-competency gap
- **Strict-narration LLM layer with post-guard** — reused for per-competency narratives (5 calls per employee instead of 1)
- **Ready-to-paste with Copy button** — reused per competency
- **Anomaly flags as chips** — reused with OUTLIER / ALIGNED / REVIEW semantics
- **Multi-source signal aggregation in the parser** — extended to tag each signal with its source system
- **Design system** (Fraunces display, burnt orange accent, Part-A-vs-Part-B editorial layout) — already specified in DESIGN.md, fits the PDF layout naturally
- **Positioning, product validation, office-hours brief** — unchanged

Recommended path forward

Ship the current MVP as the team-overview view. It is already built, tested, and demo-ready. Do not throw it out. The PDF layout sits *inside* the existing detail page as a richer replacement for the current single-narrative view.

Rough scope for the per-employee review page (post-MVP):

- **Day 11-12** — Competency model and signal mapping. New `competencies` table keyed by (employeekey, competencyname). Five rows per employee. Migrate one department's ingest to produce 5 per-competency scores. Config file mapping each signal to one or more competencies (for example, `LMS.cert_hours`)

belongs to Growth & Development; `Jira.on_time_completion_rate` belongs to Operational Excellence).

- **Day 13** — Original-review ingestion. Accept a new `xlsx` shape (or a manual form) carrying manager rating + 1-2 sentence comment + 5 employee self-assessment bullets per competency.
- **Day 14** — Build the `/review` page matching the PDF layout. Part A on the left (original review with "generic" and "unverified" flags), Part B on the right (data score + gap + verified signals + per-competency narrative). Per-competency accordion.
- **Day 15** — Per-competency narrative generation. Five Haiku calls per employee. Generic-comment detector (simple LLM classifier: "Could this comment apply to any employee in any company? Yes/No."). PDF export using the existing `reportlab` pipeline.
- **Day 16** — Three-findings summary generator for the cycle close (Manager Comment Problem, Hidden Impact Problem, Calibration Risk).

Estimated scope: roughly another week of engineering, cleanly additive to what already ships.

What stays unchanged

- Parser, joins, DB schema (extend, do not replace)
 - LLM client, guard, mock (reuse verbatim)
 - Design system
 - Positioning, product validation, office-hours brief
 - All 64 unit tests
-

Decision needed from the product reviewer

- **Does this Before-After PDF become the new north star?** If yes, Engineering captures it in the repo alongside `product-validation.md` and `office-hours-findings.md`, and scopes the per-employee review page as the next phase in `PLAN.md`.
 - **Does the current MVP still ship as the team-overview demo?** Recommendation: yes. It proves the thesis to the HR stakeholder and has already earned the engineering investment. It becomes the entry point to the per-employee review when that page is built.
 - **Which department's signal-to-competency mapping do we build first?** The PDF uses an Engineering-heavy example (Salesforce, Jira, LMS, Scribe). That matches our Engineering data shape most closely. Recommendation: Engineering first, Sales second, HR third.
-

One-line summary

The MVP and the Before-After PDF describe the same product at different zoom levels. The MVP proves the thesis at the team level; the PDF specifies the next zoom-in — per-employee, per-competency, original-vs-verified. The structural gap is the competency model; every other primitive already exists. Estimated additional effort: one week.